

Hot Dog or Not: Image Classification Using Convolutional Neural Networks

Jone Egon Steinhoff, Lukas Rasocha, Mads Prip & Petr Boska Nylander

Technical University of Denmark, Department of Applied Mathematics and Computer Science, Kgs. Lyngby, Denmark

1 Objective



Figure 1: Example of images from the training dataset with their corresponding labels.

- The main objective is to develop a model that can accurately classify images into two categories: hotdog or not hotdog

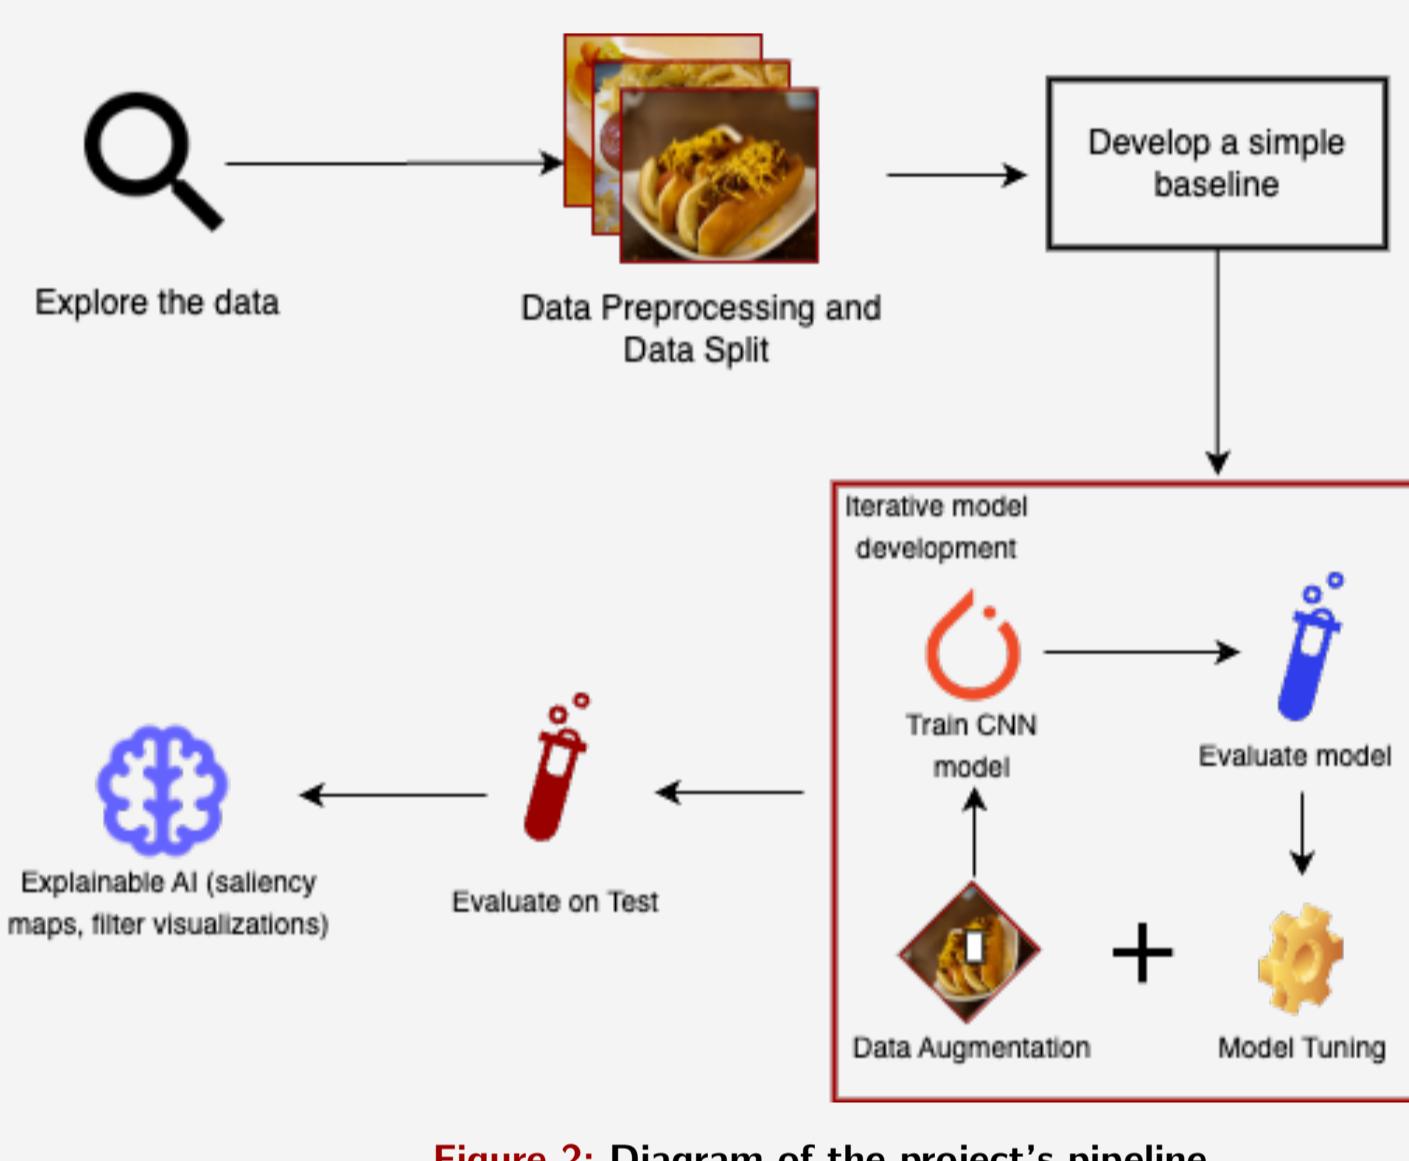


Figure 2: Diagram of the project's pipeline.

2 Data Preprocessing & Augmentation

- Training data split into validation (20%) and training (80 %)
- Data Preprocessing
 - Resize images to 224×224
 - Normalize to $\mu = 0$ and $\sigma = 1$
- Data Augmentation

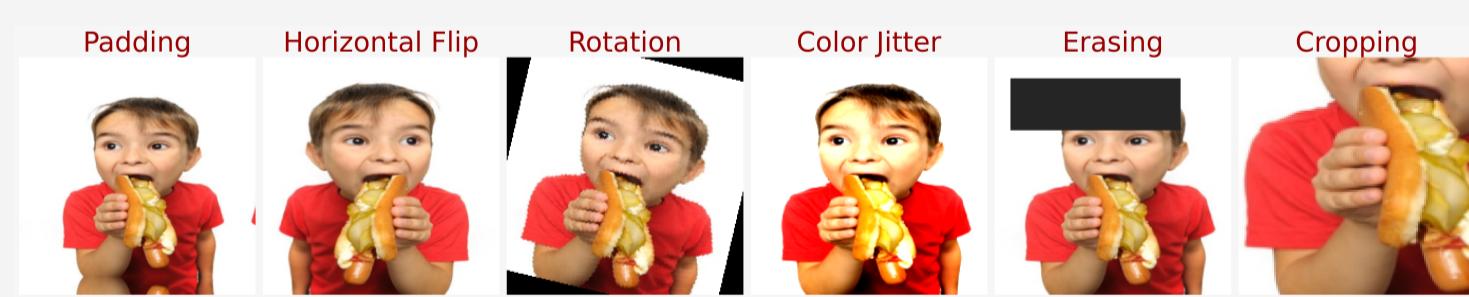


Figure 3: Augmentations applied to the input images. [2]

3 Baseline model

- Simple convolutional neural network (CNN) with 4 convolutional layers

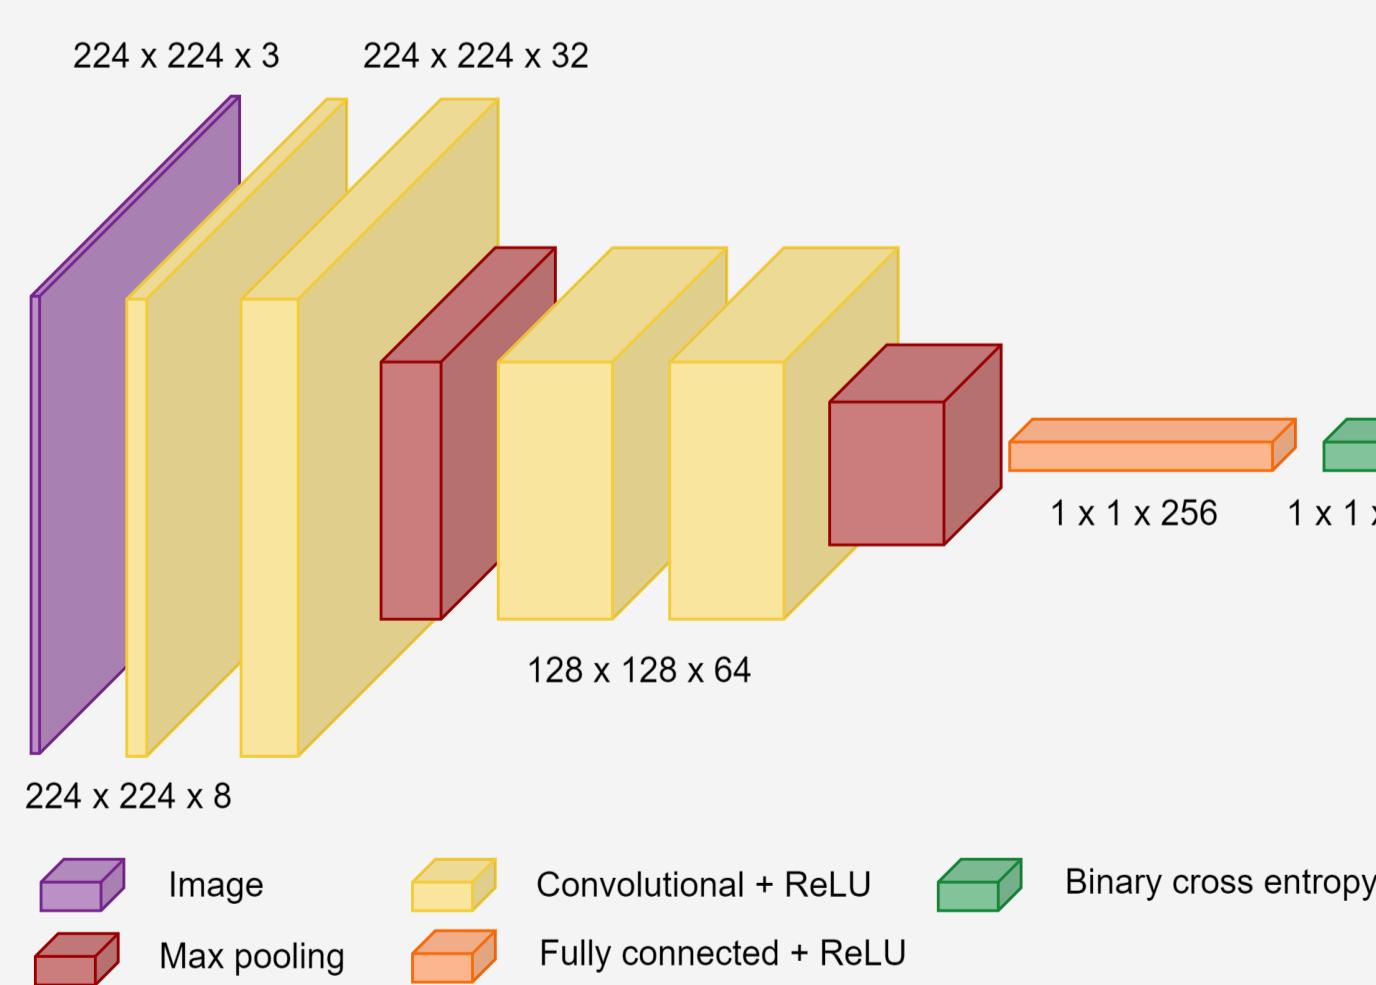


Figure 4: Architecture of the baseline model.

4 Training Process

- Reduce overfitting of the baseline CNN model

- Convolution setup:** Filter size: 3, Stride: 1, Padding: 1, Max Pooling: 2

- Training setup:** Batch size: 64, Loss Fun: Binary Cross Entropy, Epochs: Early stopping with patience 5 & Learning Rate: 0.0001 (for other details see diagram below)
- Final model result reported on a holdout test set

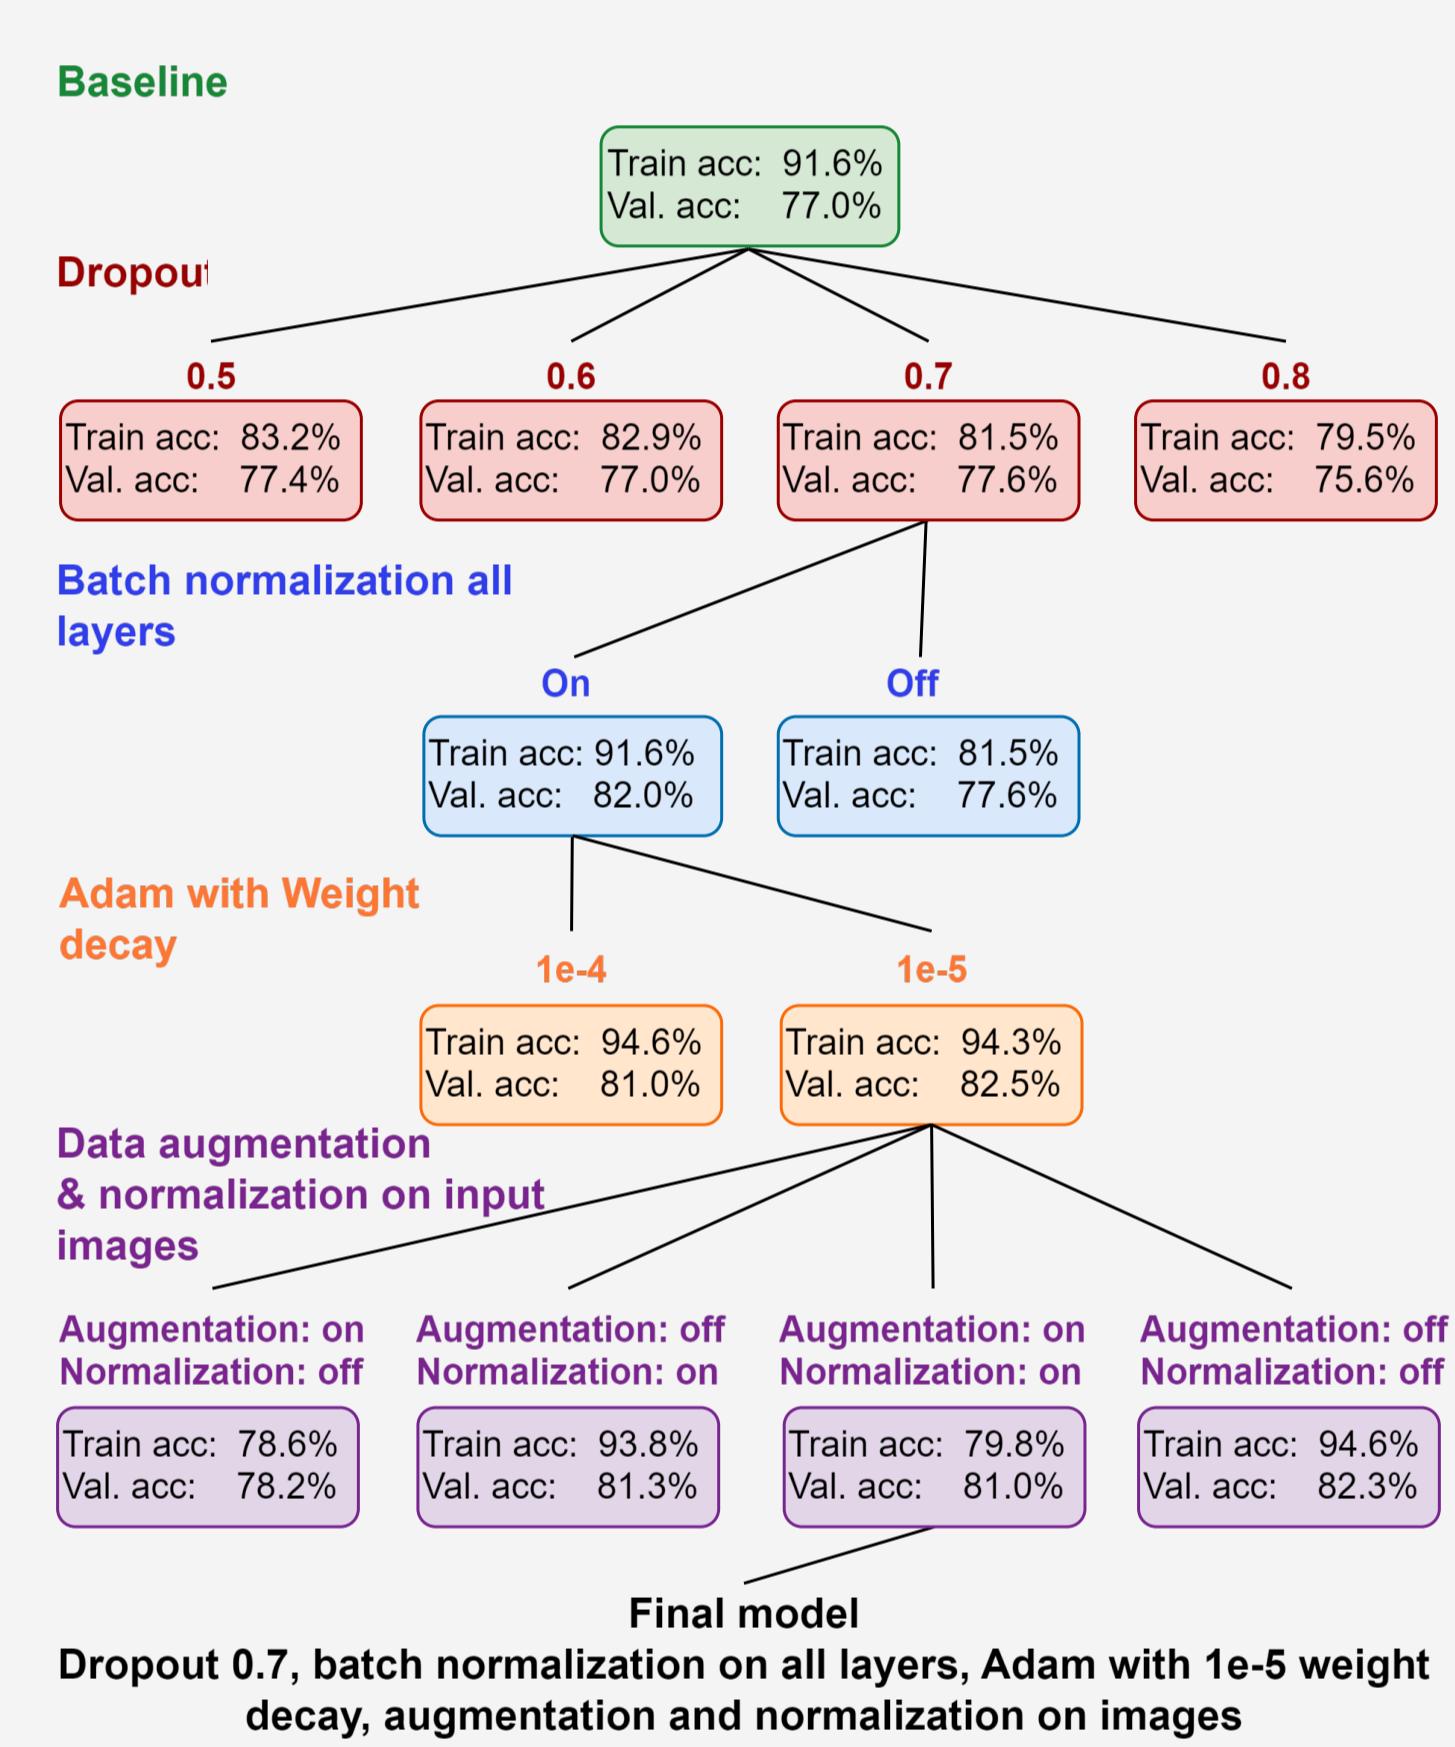


Figure 5: Parameter tuning

5 Challenges & Solutions

- Overfitting:** Implemented dropout, L2 regularization and data augmentation with early stopping
- Variance in results:** Results in each model averaged over 5 runs
- Computational resources:** Limited to not being able to use multiple combinations to find the most optimal set of hyperparameters
- Slow / bad convergence:** First lowered learning rate of SGD optimizer and later switched to Adam optimizer

6 Results

- Final obtained model:

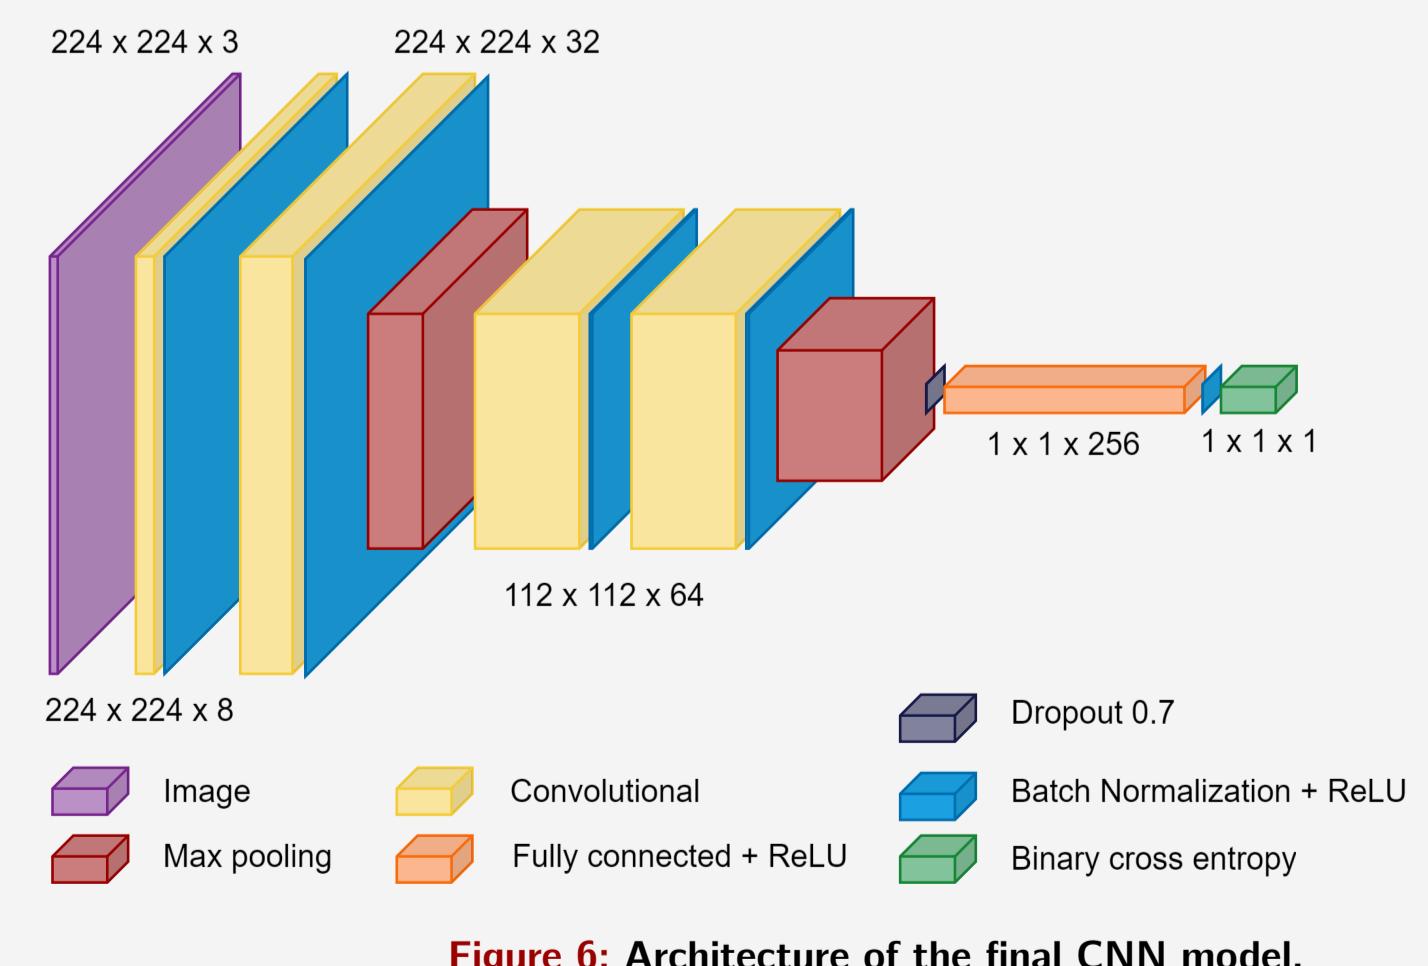


Figure 6: Architecture of the final CNN model.

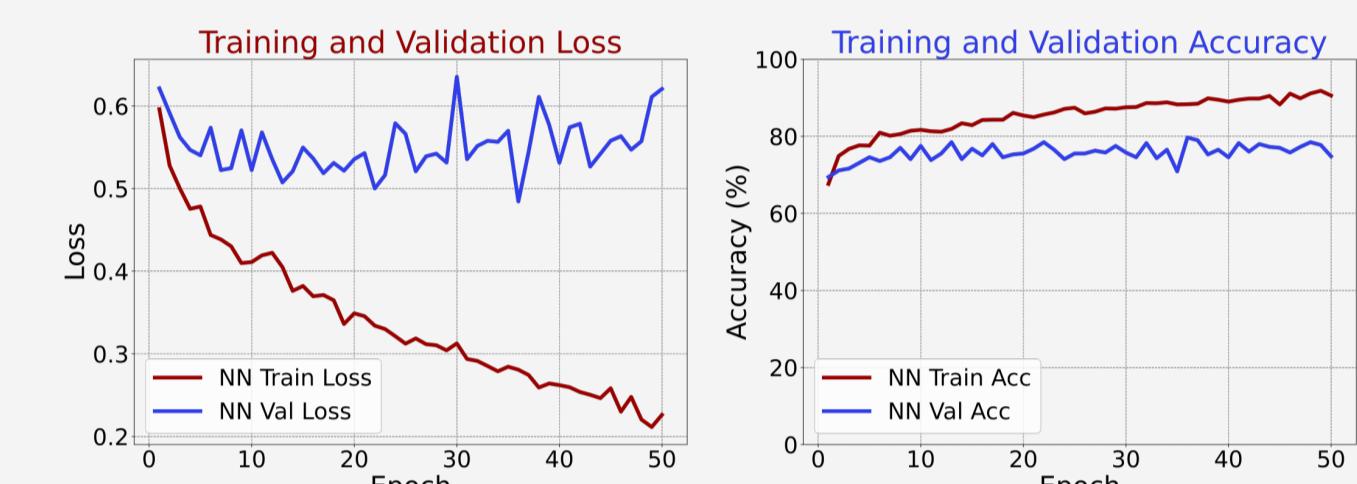


Figure 7: Training and validation loss and accuracies displayed for the baseline (top) and the final model (bottom).

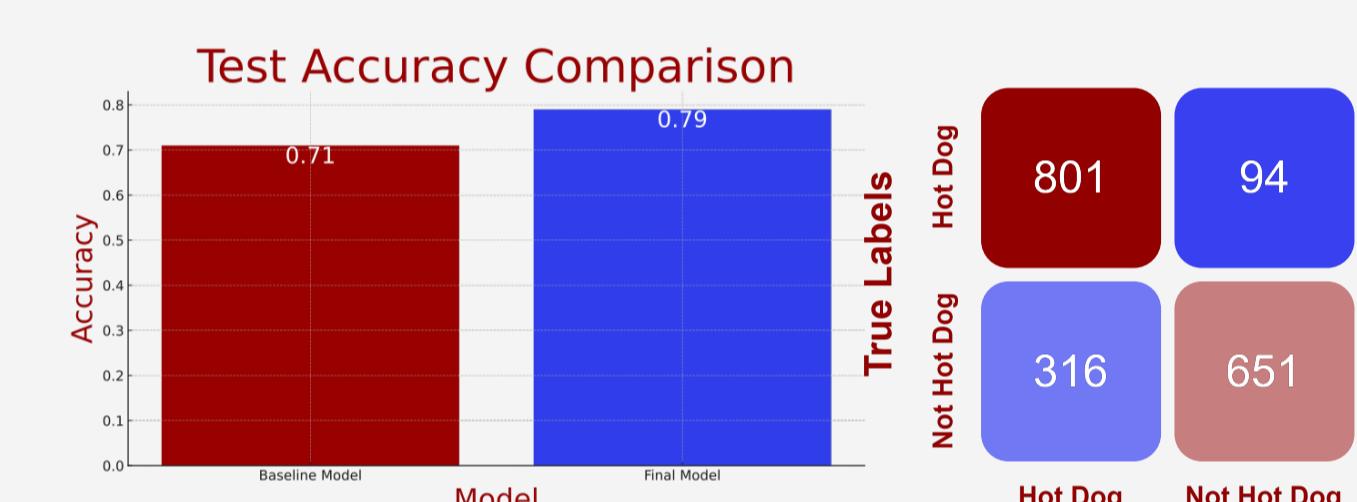


Figure 8: Test accuracy for the baseline model and the final CNN model (left). Confusion matrix on the test set for the final model (right).

7 Conclusion & Learnings

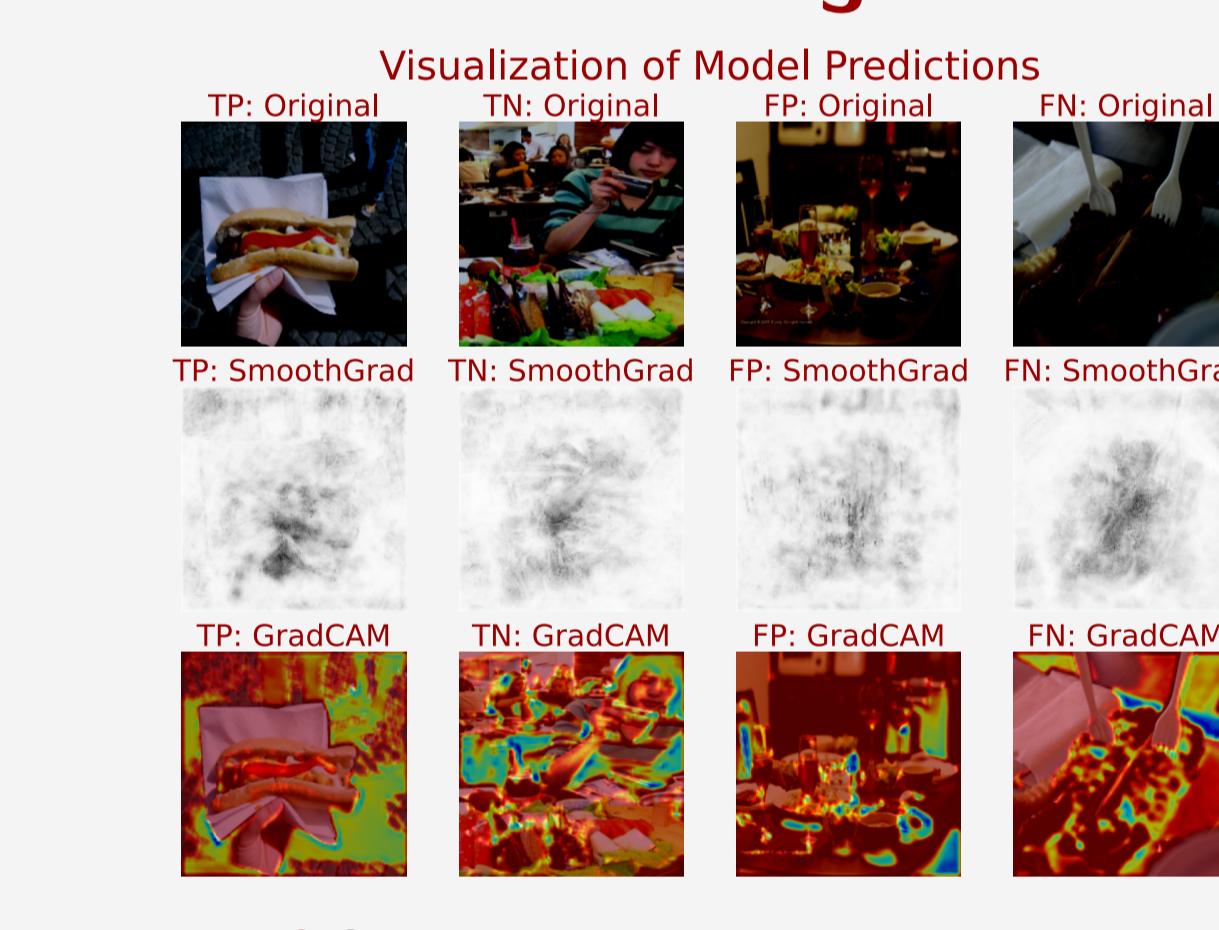


Figure 9: Saliency maps. Top row: normalized original images, middle row: SmoothGrad applied, bottom row: GradCAM applied. [1]



Figure 10: Feature maps from the first CNN layer of a random hot dog image.

- Built a CNN model with a good accuracy for binary image classification
- Learned the importance of data preprocessing and augmentation in improving model performance
- Learned how to debug and optimize deep learning models
- Future work:** Look into what categories (people, pets, furniture, food) the trained CNN model struggles to predict the most

8 References

- ChatGPT: Used for code debugging, more complicated styling of plots, explanation of complicated topics
- 1. Captum: Model Interpretability for PyTorch
- 2. Ultimate Guide to Fine-Tuning in PyTorch : Part 3. Deep Dive to PyTorch Data Transforms with Examples (Ruman, 2023)